


Block Ciphers

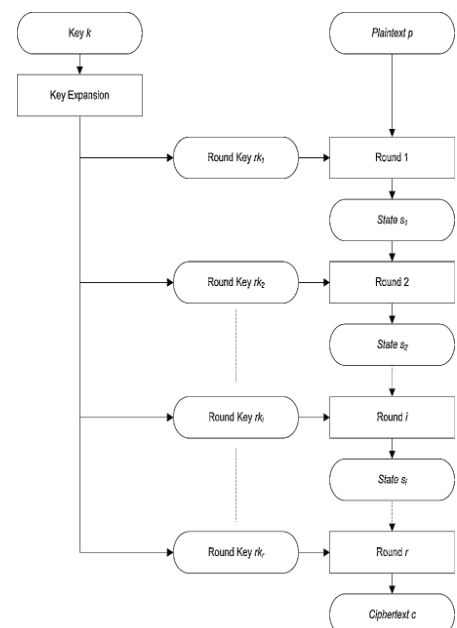
- A block cipher focuses on a simpler primitive that processes fixed-length messages (a block)



- The generalisation to a cipher proper (which deals with messages of arbitrary length) is handled by a standardised construction (**mode of operation**).
- Security-wise, the primitive is expected to be indistinguishable from a random permutation (when k is sampled randomly) — a (keyed indexed) **pseudo-random permutation (PRP)**.

Design of block ciphers

- Two distinguished desirable properties [C. Shannon 1949]:
 - Diffusion** — the influence of changes in the plaintext should be spread throughout the ciphertext;
 - Confusion** — relationship between the ciphertext and the encryption key should be as complex as possible.
- Most modern designs are iterated ciphers, where a bijective **round** transformation is applied repeatedly with distinct *round-keys* (derived from the main cipher key).



Block vs. Stream Ciphers

- Different “processing unit” (block vs. bit/byte);
- Use of block ciphers depend on some *mode of operation*;
- Stream ciphers do not promote diffusion;
- Block cipher inverse operation (decryption) is harder to obtain;
- Stream ciphers are typically faster than block ciphers;

Basic Modes of Operation

- A block cipher primitive by itself can only process fixed length bitstrings (blocks).
- A block cipher **mode of operation** describes how the primitive is used to encrypt/decrypt arbitrary messages.
- Depending on the applications, we can choose among several standard modes:
 - *Electronic Code Book (ECB)*
 - *Cipher Block Chaining (CBC)*
 - *Counter Mode (CTR)*
 - ...
- Each mode of operation comes equipped with its own security analysis: assumptions; security proofs; etc.


Padding

- Certain uses of block ciphers (modes of operation) require the size of the plaintext to be a multiple of the block size.
- This requirement is usually solved by adopting a **padding** method: a reversible transformation that fills the last block of the plaintext.
 - Several standard constructions available
 - bit vs. byte oriented;
 - deterministic vs. randomised;
 - etc.
 - E.g. PKCS7 padding


PKCS#7 Valid Padding															
A'	B'	C'													
41	42	43	05	05	05	05	05								
A'	B'	C'	D'												
41	42	43	44	04	04	04	04								
A'	B'	C'	D'	E'											
41	42	43	44	45	03	03	03								
A'	B'	C'	D'	E'	F'										
41	42	43	44	45	46	02	02								
A'	B'	C'	D'	E'	F'	G'									
41	42	43	44	45	46	47	01								
A'	B'	C'	D'	E'	F'	G'	H'								
41	42	43	44	45	46	47	48	08	08	08	08	08	08	08	08

- [remark: *Ciphertext Stealing* is an alternative strategy for handling messages of arbitrary size].

Electronic Code Book (ECB)



Electronic Codebook (ECB) mode encryption




Electronic Codebook (ECB) mode decryption


- Requires Padding.
- Repetition of blocks is detectable - **code book attack**;
- ...leading to the leakage of patterns of the plaintext.
- **It should only be used to encrypt single-block messages.**
- Vulnerable to repetition/replacement attacks.



Cipher Block Chaining (CBC)



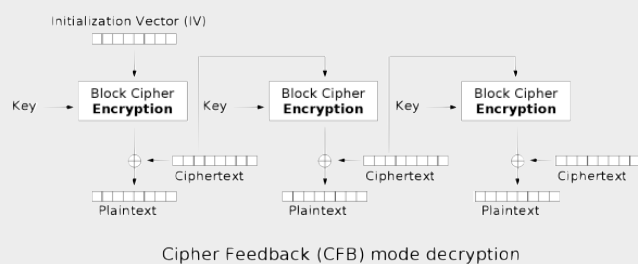
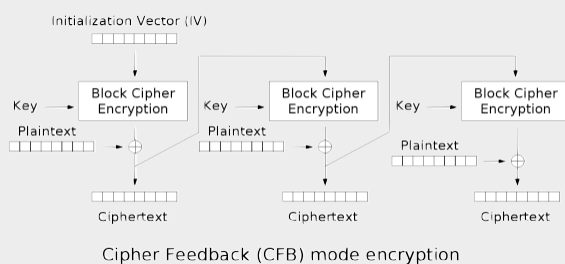
Cipher Block Chaining (CBC) mode encryption



Cipher Block Chaining (CBC) mode decryption



- Each plaintext block is masked by the previous ciphertext
- ...which induces a dependency on all previous plaintext blocks.
- A random **Initialisation Vector (IV)** shall be used to initiate the process.
- As with Nonces, the IV doesn't need to be kept secret.
- The swapping of one bit of the ciphertext would impact the corresponding block plus a single bit on the next-block.
- The last block of the ciphertext can be used as a Message Authentication Code (CBC-MAC).
 - In such use, a fixed IV shall adopted;
 - Secure for fixed-length messages (but simple fixes can be adopted to make it suitable for arbitrary length messages).

Cipher FeedBack mode (CFB)





- Implements a self-synchronising stream cipher.
- Does not requires padding.
- IV should be a Nonce.
- Notice that block cipher primitive is always used in “encrypt” mode.
- Keystream depends on the key, IV, and all the previous plaintext.
- Number of bits in the feedback is variable (CFB_n).
 - Between 1 and the block size;
 - Feedback transfers the n most significant bits to the least significant bits (with a shift of the remaining bits);
 - Number of bits in the feedback impact the resynchronisation time.

Output FeedBack mode (OFB)



- Implements a synchronous stream cipher (in output-feedback mode).
 - The block cipher is used as the *next-state function*;
 - The *output-function* is the identity function.
- Does not requires padding.
- IV should be a Nonce.
- Block cipher primitive is only used in “encrypt” mode.
- Keystream doesn’t depends on the message — is obtained by iterating the block cipher on the IV.
 - It can be expanded prior to the reception of the message;
 - but needs to be computed sequentially.
- The swapping of one bit of the ciphertext only impacts a single bit.

CouTeR mode (CTR)



- Implements a synchronous stream cipher (in counter mode).
 - The *next-state function* is a simple increment;
 - The block cipher is used as the *output-function*.
- Does not requires padding.
- IV should be a Nonce.
- Block cipher primitive is only used in “encrypt” mode.
- Keystream doesn’t depends on the message.
 - It can be expanded prior to the reception of the messsage;
 - and each block of the keystream can be processed independently (e.g. in parallel).
- The swapping of one bit of the ciphertext only impacts a single bit.

Authenticated Encryption

- A recent trend has been the adoption of block cipher modes of operation which provides both **confidentiality** and authenticity (**integrity**) of data.
- Usually supports also “*Associated Data*” — data which is not encrypted but is attached to the integrity guarantee (e.g. metadata).
- Some examples:
 - **EAX** (Encrypt-then-Mac-then-Translate)
 - **CCM** (Counter with CBC-MAC)
 - **GCM** (Galois/Counter Mode)
 - **OCB** (Offset CodeBook)

Advanced Encryption Standard (AES)

- NIST call for block cipher algorithm that would replace Data Encryption Standard (DES) — (announcement:1997, submissions: 1998, decision 2001).
- Requirements:
 - Blocks and keys with, at least, 128 bit;
 - Faster and more secure than TripleDES;
 - Detailed specification and design rational;
 - Suitable for software implementations (32bit architectures).
- Finalists:
 - **MARS** — complex, fast, high security margin
 - **RC6** — very simple, very fast, small security margin
 - **Rijndael** — clean design, fast, good security margin (**WINNER**)
 - **Serpent** — slow, clean, very high security margin
 - **Twofish** — complex, very fast, high security margin


AES (RijnDael) description

- Block size of 128 bit (16 byte), organised as a 4x4 byte matrix;
- Iterated block cipher. Number of rounds: 10, 12 or 14 (depending on the key size);
- Supports keys of 128, 192 and 256 bit. KeyExpansion algorithm expands it into required 128 bit round keys;
- Detailed mathematical justification for specification details;
- Round processing steps:
 - SubBytes
 - ShiftRows
 - MixColumns
 - AddKey(with corresponding *inverted* variants and subtle teaks to facilitate implementation)
- **Believed** to be resistant to all *known* attacks!

AES round


- ByteSub

- A single S-box
- Highly non-linear




- ShiftRow

- Rotate rows



- MixColumn

- Interaction with ShiftRow promotes high diffusion in multiple rounds



- KeyAddition

$$\begin{bmatrix} a_{0,0} & a_{0,1} & a_{0,2} & a_{0,3} \\ a_{1,0} & a_{1,1} & a_{1,2} & a_{1,3} \\ a_{2,0} & a_{2,1} & a_{2,2} & a_{2,3} \\ a_{3,0} & a_{3,1} & a_{3,2} & a_{3,3} \end{bmatrix} + \begin{bmatrix} k_{0,0} & k_{0,1} & k_{0,2} & k_{0,3} \\ k_{1,0} & k_{1,1} & k_{1,2} & k_{1,3} \\ k_{2,0} & k_{2,1} & k_{2,2} & k_{2,3} \\ k_{3,0} & k_{3,1} & k_{3,2} & k_{3,3} \end{bmatrix} = \begin{bmatrix} b_{0,0} & b_{0,1} & b_{0,2} & b_{0,3} \\ b_{1,0} & b_{1,1} & b_{1,2} & b_{1,3} \\ b_{2,0} & b_{2,1} & b_{2,2} & b_{2,3} \\ b_{3,0} & b_{3,1} & b_{3,2} & b_{3,3} \end{bmatrix}$$

One-Way Functions

One-Way Functions

- Security of cryptographic techniques often translates into establishing that a given function is **one-way**.

A function $f:A \rightarrow B$ is **one-way** when, given some element $x:A$, it is easy to compute $f(x):B$, but giving any $y:B$ in the range of f , is very difficult to find some $x:A$ such that $y=f(x)$ (a *pseudo-inverse* of y).

- Notions of "easy" and "difficult" must be understood in the context of complexity theory, depending on the existence (or not) of *polynomial-time algorithms* for their calculation.

Existence of one-way functions $\Rightarrow P \neq NP$

(obs.: but is indeed a *stronger assumption*)

- From a theoretical point of view, the existence of OW functions is usually identified as the assumption that makes symmetric cryptography feasible.

Cryptographic Hash functions

- A paradigmatic example of one-way functions are cryptographic hash functions.
- A **cryptographic hash function** is a one-way function that maps arbitrary length bit strings into a fixed-length n -bit range.

$$H : \{0,1\}^* \rightarrow \{0,1\}^n$$

- Looking at the cardinalities of domain/codomain, we clearly have that every hash function is not injective — **collisions always exist!** (that is, distinct domain points m_1 and m_2 such that $H(m_1)=H(m_2)$).
 - ...but, being OW functions, finding collisions should be difficult.
- Their use is prevalent in cryptography.

Defining properties

- The requirements of hash functions are usually expressed by the following hierarchy of properties:
 1. **(First) pre-image resistant:** given a hash value h , it should be unfeasible to obtain a message m such that $H(m)=h$.
 2. **Second pre-image resistant:** given a message m_1 , it should be infeasible to obtain a message m_2 distinct from m_1 such that $H(m_2)=H(m_1)$.
 3. **Collision resistant:** it is unfeasible to find distinct messages m_1 and m_2 such that $H(m_1)=H(m_2)$.
- The first simply asks for one-wayness of the hash function H .
- But when designing/choosing hash functions, we look at the stronger property of *collision resistance* (cannot find any collision).
- However, the security of some applications only ask for one of the a weaker properties...

Birthday paradox

- A standard result in probability theory tell us that we need reasonably sized codomain to achieve collision resistance.

How many people need to meet to make it more likely than unlikely that at least 2 share the same birthday?

- ...it would be enough to check around $\sqrt{366}$ random elements!
- If we consider typical hash codomain sizes ranging on 128 and 512 bits
- ...we conclude that the *birthday attack* would require around 2^{64} and 2^{256} random picks to find a collision.

Examples

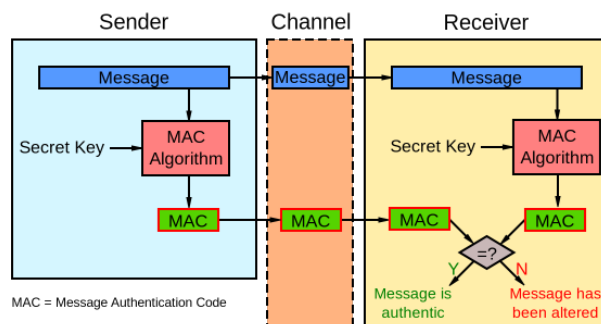
- Most common cryptographic hash functions:
 - **MD5** — codomain size: 128 bit. **Already broken!**
 - **SHA-1** (Standard Hash Algorithm) — codomain size: 160 bit. **Already broken!**
 - **SHA-2** — codomain sizes: 224, 256, 384 and 512 bits. Adopt the same design of SHA-1, with enhanced security.
 - **SHA-3** — codomain sizes: 224, 256, 384 and 512 bits. Follow a different design (sponge construction).
 - **SHAKE-128/256** — eXtendable Output Function (XOF). The same design of SHA-3.

Applications of Hash functions

- Password storage/verification;
- Building block of other cryptographic components:
 - Message Authentication Codes (MAC);
 - Key-Derivation Functions (KDF);
 - Secure Pseudo-Random Number Generators;
 - Commitment-schemes;
 - Block ciphers;
 - ...

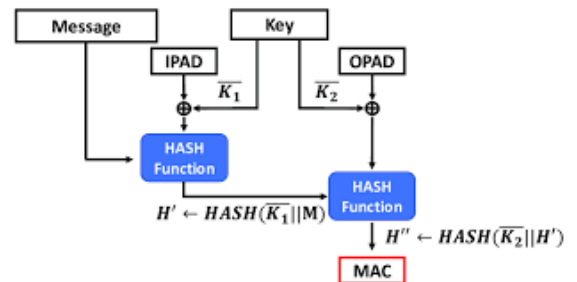
Message Authentication Codes (MAC)

- Hash functions, per se, do not ensure integrity of data — a malicious user might recompute its value for the modified data.
- A **Message Authentication Code (MAC)** should be used — it relies on a secret key that is known only to the legitimate parties (sender and receiver).
 - sender computes a **authentication tag** (from message and key);
 - receiver recomputes the tag, and compares it with the received one.
- Informally, it can be thought as a Keyed-Hash function (Hash function with a secret key).
- Examples: Poly1305; CMAC.



HMAC

- A MAC can be easily constructed from any Hash function.
- A standard construction is the Hash-based Message Authentication Code (HMAC).
 - It performs two-passes on the hash function, to account for possible weakness of the used functions (length-extension attacks).
- Examples: HMAC-SHA256; etc.
- Remark: the SHA-3 design is not vulnerable to length-extension attacks. Hence, the HMAC nested construction is redundant for the SHA-3 family of hash functions — its enough to use $H(K||M)$.



Key-Derivation Functions (KDF)

- Key-Derivation Functions (KDF) allow the derivation of cryptographic keys from other (secret) material.
- KDFs are used in a variety of situations:
 - derive keys from weak secrets (passwords/passphrases);
 - derive keys for different cryptographic purposes from a single “*master secret*” (key diversification);
 - derive keys of different length from the ones provided;
- Sometimes, it accepts also non-secret inputs (other-info), allowing to bind the secret to application-specific data.

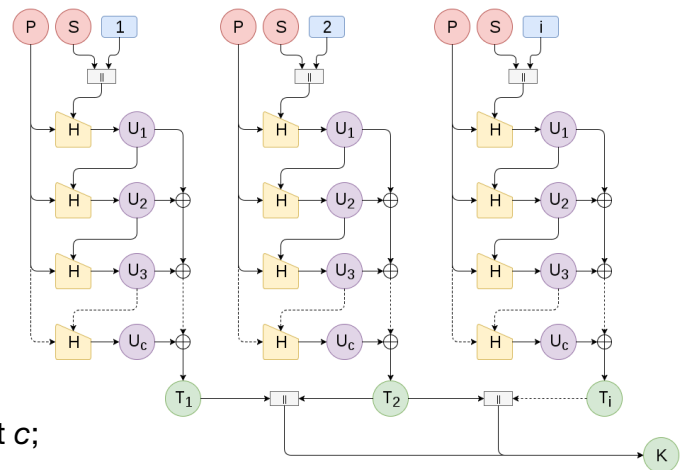
Password-Based KDF

- Passwords/Passphrases cannot be used directly as cryptographic keys
- ...but we can derive cryptographic keys from those passwords (through a KDF).
- Main problem is that a KDF doesn't solve their inherent weaknesses:
 - low entropy;
 - vulnerable to *dictionary attacks*.
- Despite all their drawbacks, passwords remain a prevalent resource in information security (user authentication; keystore protection; etc.)

Password-Based KDF

- Passwords/Passphrases cannot be used directly as cryptographic keys
- ...but we can derive cryptographic keys from those passwords (through a KDF).
- Main problem is that a KDF doesn't solve their inherent weaknesses:
 - low entropy;
 - vulnerable to *dictionary attacks*.
- Despite all their drawbacks, passwords remain a prevalent resource in information security (user authentication; keystore protection; etc.)

- It is recommended that a KDF, to be used with weak secrets, implement countermeasures that minimise the problems mentioned above:
 - use of a **salt** — a random number that randomises the output of the OW-function. It helps defending against attacks that use pre-computed tables (e.g. rainbow tables);
 - be **resource intensive** (CPU time and/or memory) to make dictionary attacks more difficult.
- Example: PBKDF2
 - Parametric on some MAC H ;
 - uses a salt S and a iteration count c ;
 - produces a secret of some given length.



Password protection

- Certain applications, and specifically *password protection*, recommend stricter counter-measures to dictionary attacks.
 - account for the possibility of attackers build dedicated password-cracking machines;
 - go beyond computational intensive: *memory-hard* designs.
- Some algorithms:
 - **bcrypt** (1999) — support adjustable cost;
 - **SCrypt** (2008) — combines PBKDF2 with a special construction (ROMix) making it resistant to hardware brute-force attacks;
 - **Argon2** (2015) — winner of the Password Hashing Competition.

Security balance

- ...but in the end, security will always depend on the entropy of the passphrase!

Estimated cost of hardware to crack a password in 1 year.

KDF	6 letters	8 letters	8 chars	10 chars	40-char text
DES CRYPT	< \$1	< \$1	< \$1	< \$1	< \$1
MD5	< \$1	< \$1	< \$1	\$1.1k	\$1
MD5 CRYPT	< \$1	< \$1	\$130	\$1.1M	\$1.4k
PBKDF2 (100 ms)	< \$1	< \$1	\$18k	\$160M	\$200k
bcrypt (95 ms)	< \$1	\$4	\$130k	\$1.2B	\$1.5M
scrypt (64 ms)	< \$1	\$150	\$4.8M	\$43B	\$52M
PBKDF2 (5.0 s)	< \$1	\$29	\$920k	\$8.3B	\$10M
bcrypt (3.0 s)	< \$1	\$130	\$4.3M	\$39B	\$47M
scrypt (3.8 s)	\$900	\$610k	\$19B	\$175T	\$210B

Key-Management

Cryptographic Keys

- A critical factor in the security of cryptographic techniques is the quality of the keys that are used.

...always use Cryptographic Secure Random-Number Generators!

- Their shape (and specifically, their size) depend on the specific technique. For symmetric cryptography, keys are usually random bit strings (sizes ranging between 128 and 256 bits).
- Extreme care must be taken when handling cryptographic keys (storage; backup; disposal; etc.)

Key handling

- As a rule of thumb:
 - the *key lifetime* (**cryptoperiod**) shall be as short as possible,
 - and inversely proportional to their use and the criticality of the protected data.
- Keys are classified as
 - **Long-term** keys (aka **static** keys);
 - **Short-term** keys (aka **ephemeral** keys).
- As we will see, symmetric cryptography favours short-term keys.
- E.g. the establishment of a secure-channel should rely on session keys:

A **session key** is a single-use symmetric key used for safeguarding communication during a specific interaction (a **session**).
- On the programming level, cryptographic APIs typically handle keys as an *abstract datatype*.
 - allow strict control on the available operations to manipulate keys;
 - support specific safeguards (e.g. clear the memory after use).

Key distribution

- The pre-distribution of keys is the major challenge in using symmetric crypto.
- They typically depend on secure-channels, which are costly.
- Notice the circularity:

Cryptography can be used to establish a secure channel between two parties, but itself depends on the existence of a secure channel to distribute the secret key.

- Remembering the recommendation to use session keys makes the problem worse!
- We will see that asymmetric cryptography offers more attractive solutions, but is instructive to consider how the use of symmetric cryptography can obviate the problem.

Key distribution protocols

- In a community with N users, the number of keys each member needs to store to allow secure communication between any pair of members is $(N * (N - 1))/2$.
 - This is clearly not feasible in a community of considerable size (e.g. internet).
- **Key Distribution Protocols** make use of a "trust network" to distribute the keys between concerned parties.
 - E.g.: a Trusted Third-Party (TTP) shares a key with each member;
 - The TTP generates and distributes a session-key between any pair of agents on demand.
- But these protocols are still not suitable for open communities (such as the Internet).

Guidelines for key-management

- Keys shall not be used for multiple purposes (use instead KDFs to derive different keys);
- In security-critical operations, key-handling (generation; storage; use; etc.) should be restricted to **Hardware Security Modules (HSM)**.
- Keys should be deleted when compromised or expired (but different applications can adopt different policies)
- If key-backup mechanism is needed, a threshold *secret-sharing scheme* should be used.

Cryptographic Key-Management Systems (CKMS)

- A CKMS consists of a set of policies, components and devices that are used to protect, manage and distribute cryptographic keys.
- A CKMS covers all states a key passes between its generation and its destruction — **key lifecycle**

